

# Information-centric Networking and Security

Edited by

Edith Ngai<sup>1</sup>, Börje Ohlman<sup>2</sup>, Gene Tsudik<sup>3</sup>, and Ersin Uzun<sup>4</sup>

<sup>1</sup> Uppsala University, SE, [edith.ngai@it.uu.se](mailto:edith.ngai@it.uu.se)

<sup>2</sup> Ericsson Research - Stockholm, SE, [borje.ohlman@ericsson.com](mailto:borje.ohlman@ericsson.com)

<sup>3</sup> University of California - Irvine, US, [gts@ics.uci.edu](mailto:gts@ics.uci.edu)

<sup>4</sup> Xerox PARC - Palo Alto, US, [ersin.uzun@acm.org](mailto:ersin.uzun@acm.org)

---

## Abstract

In recent years, Information-centric Networking (ICN) has received much attention from both academic and industry participants. ICN offers a data-centric means of inter-networking that is radically different from today's IP-based Internet which is host- or address-centric. Security and privacy issues in ICN have become increasingly important, as ICN technology gradually matures and nears real-world deployment. As is well known, in today's Internet, security and privacy features were added ad hoc and had to be incrementally and individually retrofitted (with varying success) over the last 35 years. In contrast, since ICN-based architectures (e.g., NDN, CCNx, etc.) are still evolving, it is both timely and important to explore ICN security and privacy issues as well as devise and assess possible mitigation techniques. Therefore, the general purpose of this Dagstuhl seminar was to discuss and explore potential ICN security features, attacks, privacy leaks, and potential means of mitigating vulnerabilities. During the seminar, these points were the focus of our discussion:

- What are the relevant threat models with which ICN must be concerned?
- To what extent is trust management a solved problem in ICN? Have we adequately identified the core elements of a trust model, e.g., with NDN trust schemas?
- How practical and realistic is object-based security when framed in the context of accepted privacy measures used in IP-based networks? Ephemeral keys and forward-secure sessions invalidate ICN caches but allow us to
- Are there new types of cryptographic schemes or primitives ICN architectures should be using or following that will enable more efficient or secure packet processing?

The seminar satisfied many of these questions and fueled discussions for those remaining. All participants briefly self-introduced themselves. Afterwards, select participants contributed talks on various topics of interests, ranging from trust management, privacy, to threat models. Between these primary discussions, the group partitioned itself into various working groups to spend a significant amount of time focusing on a select topic of interest to the subgroup. Example working group topics included routing on encrypted names, ICN and IoT, non-privacy-centric aspects of ICN security, and trust and identity in ICN. Following these working groups, a representative from each would present the findings to the group. (These are documented in the remainder of this report.)

The discussions at this seminar elucidated the lack of clarity the community has on topics that are important to the continuation of ICN. For instance, namespace management is still a topic of much ambiguity and confusion. There is also no consensus about the need for forward secrecy in ICN as a foundation for future networks. It is our hope that these lengthy debates serve as a fruitful resource for future research into some of the biggest security and privacy elements of ICN.



Except where otherwise noted, content of this report is licensed under a Creative Commons BY 3.0 Unported license

Information-centric Networking and Security, *Dagstuhl Reports*, Vol. 6, Issue 06, pp. 1–11

Editors: Edith Ngai, Börje Ohlman, Gene Tsudik, and Ersin Uzun



Dagstuhl Reports

Schloss Dagstuhl – Leibniz-Zentrum für Informatik, Dagstuhl Publishing, Germany

We thank Schloss Dagstuhl for the environment necessary to galvanize members of the ICN community to tackle these difficult problems. Much progress was had over the course of the seminar and since its completion, and this is primarily because of the ease of face-to-face collaboration and interaction held at Dagstuhl.

**Seminar** June 19–22, 2016 – <http://www.dagstuhl.de/16251>

**Digital Object Identifier** 10.4230/DagRep.6.6.1

## 1 Table of Contents

### Overview of Talks

Threat Models for ICN <i>Ersin Uzun</i> . . . . .	4
Cryptographic Algorithms and Security Protocols for ICN <i>Christopher A. Wood</i> . . . . .	4
Violating Consumer Anonymity: Geo-locating Nodes in Named Data Networking <i>Mauro Conti</i> . . . . .	4
ABE in ICN <i>Ashish Gehani</i> . . . . .	5
PROJECT ORIGIN: A peer-to-peer object store for CCN <i>Marc Mosko</i> . . . . .	5

### Working groups

Names and Locators <i>Marc Mosko</i> . . . . .	6
ICN and IoT <i>Edith Ngai</i> . . . . .	6
Security, Not Privacy <i>Craig Partridge and Ghassan Karame</i> . . . . .	7
Trust and Identity in ICN <i>Jan Seedorf, Kenneth L. Calvert, and Christopher A. Wood</i> . . . . .	7
Routing on Encrypted Names <i>Christopher A. Wood, Edith Ngai, Jan Seedorf, and Matthias Wählisch</i> . . . . .	8


### Open problems

Revisiting "Securing the Data Not The Pipe" <i>Marc Mosko</i> . . . . .	9
User-Generated Content in the FIB <i>Thomas C. Schmidt</i> . . . . .	9
Motivating Transport Privacy for Data Structures <i>Christian Tschudin</i> . . . . .	10
Whither ICN Privacy? <i>Gene Tsudik</i> . . . . .	10
Forward Security in ICN? <i>Christopher A. Wood</i> . . . . .	10

## 2 Overview of Talks

### 2.1 Threat Models for ICN


*Ersin Uzun (Xerox PARC - Palo Alto, US)*

License  Creative Commons BY 3.0 Unported license  
© Ersin Uzun

Networking helped to create today's world of content but was not designed for it. ICN attempts to move away from: (a) the communication model that is all about hosts and a point-to-point conversation between them, (b) the host-based central abstraction in the network, and (c) security problems of the current IP-based Internet architecture. The ICN emphasis on object security instead channel security is one step towards this transition. To quantify the degree by which security is improved, a threat model is needed. These threat models must be tailored toward the particular design challenge, such as infrastructure protection, user-friendly key distribution and trust management and enforcement, and content protection and access control. We discuss some necessary elements of threat models for these scenarios and suggest a strategy to use them in the design process.

### 2.2 Cryptographic Algorithms and Security Protocols for ICN

*Christopher A. Wood (University of California - Irvine, US)*

License  Creative Commons BY 3.0 Unported license  
© Christopher A. Wood

Old and new cryptographic algorithms and security protocols with relevance to ICN are discussed. Topics include integrity and authenticity, privacy, and availability. The goal is to stimulate useful discussions about the security services and properties that future ICN architectures could provide.

### 2.3 Violating Consumer Anonymity: Geo-locating Nodes in Named Data Networking

*Mauro Conti (University of Padova, IT)*

License  Creative Commons BY 3.0 Unported license  
© Mauro Conti

Talking about ICN privacy, we should not focus only on What (i.e., the content of a message) and Who (i.e., the sender, or the receiver, of a message): we should also be worried about protecting "Where", meaning the geographical position (or just the position within a network topology) where the message is originated from, or destined to. In fact (as shown in the ACNS '15 paper by Compagno et al. [1]), ICN caching and interest-collapse mechanisms make ICN itself inherently vulnerable to the possibility for an adversary to locate consumers. Interestingly, an approach similar to the one to violate consumer (location) privacy, might be used also to detect eavesdropper. Discussion questions:

- Exactly what routing information is made available to routing-aware adversaries, and how useful is it to run attacks?

- By observing past traffic, can I infer where interests will be routed in the network?
- What capabilities does the adversary possess?
- How this compare to similar attacks in IP?

## References

- [1] Alberto Compagno, Mauro Conti, Paolo Gasti, Luigi Vincenzo Mancini, Gene Tsudik: Violating Consumer Anonymity: Geo-Locating Nodes in Named Data Networking. ACNS 2015: 243-262.

## 2.4 ABE in ICN

*Ashish Gehani (SRI - Menlo Park, US)*

License © Creative Commons BY 3.0 Unported license  
© Ashish Gehani

The ENCODERS information-centric networking system was designed for publish-subscribe applications running over a mobile ad hoc network. It uses multi-authority attribute-based encryption to allow content access to be scoped to selected nodes in the system. Since the system is completely decentralized, peers serve as brokers that match content from publishers with interests expressed by subscribers. In order to perform such a match, an intermediate node must be authorized to see both the relevant content tags and subscriber interests. All of this has been described in a previous publication [1].

The talk described the following unpublished observation. The access control policies that applied to the metadata (content tags and subscriber interests) effectively creates reachability constraints that are independent from the one defined by the routing protocols. (This is because the access control defines which brokers can match and therefore route traffic.) Consequently, this security-routing interaction must be treated carefully during policy definition. We examined this empirically for a number of routing and access control policies.

## References

- 1 Mariana Raykova, Hasnain Lakhani, Hasanat Kazmi, and Ashish Gehani, Decentralized Authorization and Privacy-Enhanced Routing for Information-Centric Networks, 31st Annual Computer Security Applications Conference (ACSAC), 2015.

## 2.5 PROJECT ORIGIN: A peer-to-peer object store for CCN

*Marc Mosko (Xerox PARC - Palo Alto, US)*


License © Creative Commons BY 3.0 Unported license  
© Marc Mosko

Project Origin is a bittorrent-like system for CCNx that uses concepts of nameless objects and custodian routing to achieve a scalable, secure content distribution system. In this talk we present the initial design of Project Origin in the context of CCNx.

### 3 Working groups

#### 3.1 Names and Locators

*Marc Mosko (Xerox PARC - Palo Alto, US)*

License  Creative Commons BY 3.0 Unported license  
© Marc Mosko

The CCNxTorrent working group discussed locators and fetching data with non-topological names (or even topological names that are cached off-path). Routing should, it seems, only concern itself with topological names or addresses. Finding data (objects) with non-topological names should not be done in the data plane. It should be done via a service.

In CCNx, the service could resolve a named root manifest to then resolve locator names by hash. In NDN, it resolves the link routing hints to allow off-path interest forwarding. In TagNet, there is a distinction between Locator names and Descriptor names. Locator names have a strong binding between their name and a point of attachment. Descriptor (hash) names, on the other hand, are free-form and could be present anywhere. One resolves a tag query (of either type) to a topological locator and then does data transfer on that locator.

This led to a discussion on locator and identifier split. Should CCNx embrace this, or continue on with its mixed use of the name? For example, if there is a clear locator field and then a clear identifier tuple {name, [keyed restriction], [hash restriction]}, one would get full matching expressivity with the functionality of nameless object locators. A similar approach could be done in NDN, though with a different tuple. There was no consensus on this idea, though it is worth exploring.

There was some discussion on the benefit of ICN if one still needs to do an external name to address lookup. Why bother if one still needs an IP/DNS like function? One partial answer is that in the non-global routing space (i.e., data center, maybe IoT, some internal applications), one could inject all names into the internal routing protocol and realize the full benefit of application-specific names. Another argument is that it improves how one writes software to not have to deal with IP addresses and host-based networking. One could also see benefits from a re-thought network stack beyond sockets.

#### 3.2 ICN and IoT

*Edith Ngai (Uppsala University, SE)*

License  Creative Commons BY 3.0 Unported license  
© Edith Ngai

The Internet of Things (IoT) is connecting billions of smart devices (e.g. sensors) and is growing very fast. We expect more than 1 million networked “things” per square kilometer in 2030. In this group, we tried to explore how much data density we can afford and how we communicate with the “things” (say, directly to the sensors, or indirectly through the cloud or gateways). We discussed the potential of implementing ICN for the IoT. For instance, the ICN routers connecting to sensors can cache sensor data to improve the performance of data dissemination. Users can obtain data directly from the sensors and the ICN routers, without going through the cloud. There are several security concerns:

- How are sensors securely configured at the time of initialization?
- How can software updates be performed securely?

- How can we handle ICN mobility for IoT? For example, each sensor may have unique publisher identity. How do mobility and naming affect the scalability of routing?

We discussed the potentials and concerns of caching data at the sensors:

- Sensors are resource limited devices. However, memory resources may increase and the price will go down.
- It is advantageous to retrieve data directly from the sensors in some use cases (e.g., to control home lighting without going through the cloud).
- When using cryptography on sensors, the encryption time could be long and cause a delay in data retrieval.
- Sensors have to be always on to listen to the interests, which may consume a lot of energy. Scheduling or adaptive duty cycles might be considered to mitigate this.

Our summary and future plan is as follows. First, we plan to come up with sample IoT use cases, which allow us to understand more about the security needs and communication patterns in ICN for IoT. Second, we will aim at answering the following questions:

- How does IoT benefit from ICN?
- How to configure sensors at bootstrapping?
- Explore the cost of providing security for IoT data

### 3.3 Security, Not Privacy

*Craig Partridge (BBN Technologies - Cambridge, US) and Ghassan Karame (NEC Laboratories Europe - Heidelberg, DE)*

**License** © Creative Commons BY 3.0 Unported license  
© Craig Partridge and Ghassan Karame

The “No Privacy” security working group sought to answer the following question: is an ICN security architecture easier to devise if the designs fundamentally make privacy hard to achieve? In particular, the group discussed:

- What ICN entities (content consumers, hosts, routers, content creators) need identities?
- What entities can simply operate with a public/private key pair but no formal name?
- Does splitting routing out as an application help?
- Do interests need to be authenticated at each router?

The group achieved a simple security model. Members of the group hope to write up the result as a short paper.

### 3.4 Trust and Identity in ICN

*Jan Seedorf (NEC Laboratories Europe - Heidelberg, DE & Hochschule für Technik - Stuttgart, DE), Kenneth L. Calvert (University of Kentucky - Lexington, US), and Christopher A. Wood (University of California - Irvine, US)*

**License** © Creative Commons BY 3.0 Unported license  
© Jan Seedorf, Kenneth L. Calvert, and Christopher A. Wood

In an ideal ICN architecture, applications should be able to express their trust preferences or policies and let the “middleware” enforce them. This raises two important questions: (1)

what is the minimum set of policies that can be factored out of all trust models, and (2) what is the middleware that does this enforcement? The trust schemas pioneered by the NDN architecture [1] are exemplary of the core rules that might encompass a trust model. They specify what keys are allowed to sign what data. Since both keys and data are named resources in NDN and other ICN architectures, this means that a schema is, at its core, a simple name-to-name mapping that allows for arbitrary hierarchical trust models. It remains to be seen if other non-hierarchical trust models will be applicable to ICN.

To address the second question, we had to agree upon what the network is responsible for enforcing. First and foremost, network layer “trust enforcement” should not prohibit or prevent other application-layer trust models. This means that the network functionality must be simpler than that which is supported by the middleware. Currently, this is comprised of digital signature and content object hash verification. Behaviors such as certificate chain resolution or key retrieval should not be part of the core network functionality. This means that in the general “network,” routers are only responsible for single signature or hash verification. All other network nodes (e.g., consumers and producers) contain the middleware responsible for handling the remaining trust enforcement steps.

After addressing trust, we turned to identity. We discussed the following major questions:

- How are names registered and managed in ICN?
- How can names possibly be location agnostic? Is there always a discovery or locator service?

Namespace ownership is intrinsically tied to identity. Thus, namespace advertisements under different namespaces or in different networks must be authenticated with respect to the claimed owner’s identity. In this context, an identity is a public and private key pair. We struggled with questions about namespace scale and the practicality of a global namespace. Questions such as, “how do NATs work in a global namespace,” drove the discussion. No consensus or common understanding was reached.

## References

Yingdi Yu, kc claffy, Alexander Afanasyev, Van Jacobson, David Clark, and Lixia Zhang. “Schematizing trust in named data networking.” Proceedings of the 2nd International Conference on Information-Centric Networking. ACM, 2015.

## 3.5 Routing on Encrypted Names

*Christopher A. Wood (University of California - Irvine, US), Edith Ngai (Uppsala University, SE), Jan Seedorf (NEC Laboratories Europe - Heidelberg, DE & Hochschule für Technik - Stuttgart, DE), and Matthias Wählisch (FU Berlin, DE)*

**License** © Creative Commons BY 3.0 Unported license  
© Christopher A. Wood, Edith Ngai, Jan Seedorf, and Matthias Wählisch

This group started with a discussion about routing on encrypted names but ended up being a discussion about name privacy and the necessary conditions for it to be possible in different ICN architectures. In this context, we defined name privacy to be a property where a so-called “network name,” i.e., the name encoded in a packet, has no correlation or connection with the referenced content object. Specifically, name privacy means that a network name reveals nothing about the data inside the content object. Names should reveal no more than what is currently revealed by an IP address and port. After settling on this definition, we laid out our assumptions, including:



- There is no name discovery process or search engine.
- Content may be requested by an identifier (ID) such as its cryptographic hash digest. Moreover, revealing the content ID does not compromise privacy.
- Consumers know the public key of the producer to which they want to communicate.
- Network names have an implicit routable prefix and application-specific suffix. By default, consumers do not know the locator and identifier split in a name.
- Requests may specify the ID of a signature verification key or the expected content.

There are fundamentally two ways to request content: (1) with and (2) without a content ID. In case (1), a request name needs to only contain a routable prefix that will bring the request to some cache or producer which can return the corresponding content. These locators can be completely separate from the desired content and, therefore, this approach can satisfy our name privacy goal. However, without implicit knowledge about the locator for some desired content, an upper-layer service is necessary to obtain said information.

In case (2), the application-specific suffix of a name must not reveal anything about the data. To achieve this, it must be encrypted. Name encryption introduces a number of other questions, such as how to obtain the routable prefix, what key to use for encryption, and how to “protect” the result. For the base case, assume that the routable prefix is known and that the producer public key is used for name suffix encryption. If the resultant content payload is not encrypted then one may be able to infer the name from its contents. Therefore, the content response itself must also be encrypted. This requires a consumer-generated key and CCA-secure encryption since, otherwise, eavesdroppers can replay requests or perform the same generation for the target to get a decrypted response.

In all cases, we concluded that to achieve name privacy then one needs some upper-layer service. Whether its role is to provide the routable prefix for a name, encrypt the response, or to separate a content ID and locator via some other means is an orthogonal discussion. Thus, our conjecture is that name privacy is not possible natively in the network.

## 4 Open problems

### 4.1 Revisiting "Securing the Data Not The Pipe"


*Marc Mosko (Xerox PARC - Palo Alto, US)*

**License** © Creative Commons BY 3.0 Unported license  
© Marc Mosko

We revisit the question of securing the data, not the pipe. This has been the running mantra for security in information-centric networks. Is forward secrecy possible in NDN and CCN group access control by encryption? Is the "take what you want" model of group access control, dependent on long-lived keys, realistic for future networks? Is it desirable? Also, is there any role for perimeter security or is encryption enough? In this talk, we pose these questions and others to the group to stimulate a wider discussion.

## 4.2 User-Generated Content in the FIB


*Thomas C. Schmidt (HAW - Hamburg, DE)*

License  Creative Commons BY 3.0 Unported license  
© Thomas C. Schmidt

ICN names are user-generated content in FIBs. In effect, FIBs serve as (globally) replicated name set wherein any name owner can write into the set. The complexity of this state is influenced by the fact that prefix owners can always de-aggregate and create arbitrary names, even if prefixes are restrictively assigned. However, this raises questions of resource exhaustion attacks on FIBs and general complexity attacks (e.g., hash collisions). Newer attacks involve leaking information from the FIB contents to attack the forwarding plane, among others. This talk outlines the severity of these problems as an avenue to discussing potential solutions.

## 4.3 Motivating Transport Privacy for Data Structures

*Christian Tschudin (Universität Basel, CH)*

License  Creative Commons BY 3.0 Unported license  
© Christian Tschudin

Beyond ICN packets, ICN programmers will generate linked data, e.g., FLIC. We should provide techniques, services, recipes to them that make transport privacy a non-brainer.

## 4.4 Whither ICN Privacy?


*Gene Tsudik (University of California - Irvine, US)*

License  Creative Commons BY 3.0 Unported license  
© Gene Tsudik

What is the future of privacy for ICN? To what, or whom, is ICN privacy related? Existing architectures leak a significant amount of information by default, including: who requests information, whose information is requested, when content is requested, and other miscellaneous properties of information, e.g., data contents, name, size, etc. As of yet, we have not adequately addressed these privacy problems yet.

## 4.5 Forward Security in ICN?

*Christopher A. Wood (University of California - Irvine, US)*

License  Creative Commons BY 3.0 Unported license  
© Christopher A. Wood

Forward secrecy is the property that exposure of a principal's long-term secret keys does not compromise the secrecy of their previously used ephemeral (session) keys. This is a useful property to have in the presence of eavesdropping attackers intercepting and logging traffic. Specifically, it minimizes data and key compromise windows and therefore reduces the

overall number of attack vectors. However, it requires protocols and techniques for deriving ephemeral keys and then updating keys regularly. The single request-response model does not lend itself to the establishment of forward-secure security contexts without involving more esoteric cryptographic schemes. Consequently, the majority of work on ICN object security has ignored this property, which puts ICN at odds with best practice techniques used in IP-based protocols. In this talk, we seek to raise awareness of this issue and seek answers to the following important questions. First, under what conditions does transport security require forward secrecy, if at all? Second, can object encryption subsume transport security? And lastly, is forward secrecy in ICN needed?

